

CIFAR-WGAN

December 4, 2018

```
In [1]: use_gpu = True
        data_dir = '/home/kartik/data'

In [2]: import sys
        sys.path.append('../')

        %load_ext autoreload
        %autoreload 1
        %import alphagan

In [3]: from collections import defaultdict
        from psutil import cpu_count

        import numpy as np
        import pandas as pd

        import torch
        from torch import nn
        from torch.nn import init, Parameter
        import torch.nn.functional as F
        from torch.autograd import Variable
        from torch.utils.data import DataLoader
        from torchvision import datasets, models, transforms
        from torchvision.utils import make_grid

        import matplotlib.pyplot as plt
        %matplotlib inline

In [4]: # get a image data set into a tensor, scale to (-1, 1), discarding labels
        def torchvision_dataset(dset_fn, train=True):
            dset = dset_fn(
                data_dir,
                train=train,
                transform=transforms.Compose([
                    transforms.ToTensor()
                ]),
                target_transform=None,
                download=True)
            return torch.stack(list(zip(*dset))[0])*2-1
```

```

In [5]: cifar = torchvision_dataset(datasets.CIFAR10, train=True)
        cifar_test = torchvision_dataset(datasets.CIFAR10, train=False)
        print(cifar.size())

Files already downloaded and verified
Files already downloaded and verified
torch.Size([50000, 3, 32, 32])

In [6]: batch_size = 256

In [7]: num_workers = cpu_count() if use_gpu else 0

X_train = DataLoader(cifar, batch_size=batch_size, shuffle=True,
                      num_workers=num_workers, pin_memory=use_gpu)
X_test = DataLoader(cifar_test, batch_size=batch_size, shuffle=False,
                    num_workers=num_workers, pin_memory=use_gpu)

In [8]: class ChannelsToLinear(nn.Linear):
        """Flatten a Variable to 2d and apply Linear layer"""
        def forward(self, x):
            b = x.size(0)
            return super().forward(x.view(b,-1))

class LinearToChannels2d(nn.Linear):
    """Reshape 2d Variable to 4d after Linear layer"""
    def __init__(self, m, n, w=1, h=None, **kw):
        h = h or w
        super().__init__(m, n*w*h, **kw)
        self.w = w
        self.h = h
    def forward(self, x):
        b = x.size(0)
        return super().forward(x).view(b, -1, self.w, self.h)

class ResBlock(nn.Module):
    """Simple ResNet block"""
    def __init__(self, c,
                 activation=nn.LeakyReLU, norm=nn.BatchNorm2d,
                 init_gain=1, groups=1):
        super().__init__()
        self.a1 = activation()
        self.a2 = activation()
        self.norm1 = norm and norm(c)
        self.norm2 = norm and norm(c)

        to_init = []
        self.conv1 = nn.Conv2d(
            c, c, 3, 1, 1, bias=bool(norm), groups=groups)

```

```

to_init.append(self.conv1.weight)
self.conv2 = nn.Conv2d(
    c, c, 3, 1, 1, bias=bool(norm), groups=groups)
to_init.append(self.conv2.weight)

# if using grouping, add a 1x1 convolution to each conv layer
if groups!=1:
    self.conv1 = nn.Sequential(
        self.conv1, nn.Conv2d(c,c,1,bias=bool(norm)))
    self.conv2 = nn.Sequential(
        self.conv2, nn.Conv2d(c,c,1,bias=bool(norm)))
    to_init.extend([self.conv1[1].weight, self.conv2[1].weight])

# init
for w in to_init:
    init.xavier_normal(w, init_gain)

def forward(self, x):
    y = self.conv1(x)
    if self.norm1:
        y = self.norm1(y)
    y = self.a1(y)

    y = self.conv2(y)
    if self.norm2:
        y = self.norm2(y)

    return self.a2(x+y)

```

In [9]: latent_dim = 128

In [10]: *# encoder network*

```

h = 128
resample = nn.AvgPool2d
norm = None
a, g = nn.ReLU, init.calculate_gain('relu')
groups = 1#h//8
E = nn.Sequential(
    nn.Conv2d(3,h,5,1,2), resample(2), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    ChannelsToLinear(h*16, latent_dim)
)
for layer in (0,8):
    init.xavier_normal(E[layer].weight, g)

t = Variable(torch.randn(batch_size,3,32,32))
assert E(t).size() == (batch_size,latent_dim)

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:15: UserWarning: nn.init.xavier_n
from ipykernel import kernelapp as app

```

In [11]: *# generator network*

```

h = 128
norm = None
a, g = nn.ReLU, init.calculate_gain('relu')
groups = 1#h//8
resample = lambda x: nn.Upsample(scale_factor=x)
G = nn.Sequential(
    LinearToChannels2d(latent_dim,h,4,4), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    nn.Conv2d(h, 3, 1), nn.Tanh()
)
for layer in (0,9):
    init.xavier_normal(G[layer].weight, g)

t = Variable(torch.randn(batch_size,latent_dim))
assert G(t).size() == (batch_size,3,32,32)

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:16: UserWarning: nn.init.xavier_n
app.launch_new_instance()

```

In [12]: *# discriminator network*

```

h = 128
resample = nn.AvgPool2d
norm = None
a, g = lambda: nn.LeakyReLU(.2), init.calculate_gain('leaky_relu', .2)
groups = 1

D = nn.Sequential(
    nn.Conv2d(3,h,5,1,2), resample(2), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    ChannelsToLinear(h*16, 1)
)
for layer in (0,8):
    init.xavier_normal(D[layer].weight, g)

t = Variable(torch.randn(batch_size,3,32,32))
assert D(t).size() == (batch_size,1)

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:16: UserWarning: nn.init.xavier_n
    app.launch_new_instance()

```

```

In [13]: # code discriminator network
        h = 700
        a, g = lambda: nn.LeakyReLU(.2), init.calculate_gain('leaky_relu', .2)
        C = nn.Sequential(
            nn.Linear(latent_dim, h), a(),
            nn.Linear(h, h), a(),
            nn.Linear(h, 1),
        )

        for i, layer in enumerate(C):
            if i%2==0:
                init.xavier_normal(layer.weight, g)

        t = Variable(torch.randn(batch_size, latent_dim))
        assert C(t).size() == (batch_size, 1)

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:12: UserWarning: nn.init.xavier_n
    if sys.path[0] == '':

```

```

In [14]: model = alphagan.AlphaWGAN(E, G, D, C, latent_dim, lambd=40, z_lambd=0)
        if use_gpu:
            model = model.cuda()

```

```

In [15]: # model.load_state_dict(torch.load('../models/wgan_cifar10_dim_{}_lambda_{}_zlambda_{}_
        #             latent_dim, model.lambd, model.z_lambd, 4
        #             ))

```

```

In [16]: diag = []
        def log_fn(d):
            d = pd.DataFrame(d)
            diag.append(d)
            print(d)

```

```

In [17]: def checkpoint_fn(model, epoch):
        path = '../models/wgan_cifar10_dim_{}_lambda_{}_zlambda_{}_epochs_{}.torch'.format
            model.latent_dim, model.lambd, model.z_lambd, epoch
        )
        torch.save(model.state_dict(), path)

```

```

In [18]: # %%prun
        model.fit(
            X_train, X_test,
            n_iter=(1,5,5), n_epochs=4,

```

```

#     n_batches = (4,4),
    opt_params={'lr':1e-4, 'betas':(.5,.9)},
    log_fn=log_fn, log_every=1,
    checkpoint_fn=checkpoint_fn, checkpoint_every=1
)

```

```

HBox(children=(IntProgress(value=0, description='epoch', max=4, style=ProgressStyle(description=

```

```

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

```

```

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

```

	train	valid
C_critic_loss	-11.635004	-10.350653
C_gradient_penalty	1.797975	1.493553
D_critic_loss	-7.661589	-8.402270
D_gradient_penalty	1.372621	1.034688
adversarial_loss	12.419829	-23.905209
code_adversarial_loss	-1.995974	-2.857663
reconstruction_loss	12.614973	10.779352

```

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

```

```

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

```

	train	valid
C_critic_loss	-7.835329	-5.764362
C_gradient_penalty	0.969838	0.641884
D_critic_loss	-6.740700	-6.349144
D_gradient_penalty	0.869746	0.800640
adversarial_loss	-17.107382	-13.858475
code_adversarial_loss	-3.019807	-2.708389
reconstruction_loss	11.212499	11.201015

```

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

```

```

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

```

	train	valid
C_critic_loss	-4.866635	-4.289764


C_gradient_penalty	0.467750	0.383844
D_critic_loss	-4.816281	-4.687969
D_gradient_penalty	0.509644	0.503710
adversarial_loss	-3.229794	-2.691923
code_adversarial_loss	-1.526694	-0.972477
reconstruction_loss	10.788114	10.287782

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

	train	valid
C_critic_loss	-3.872909	-3.695042
C_gradient_penalty	0.341297	0.309088
D_critic_loss	-3.967117	-3.753757
D_gradient_penalty	0.380238	0.389384
adversarial_loss	-0.667429	-0.309710
code_adversarial_loss	-0.635379	-0.464543
reconstruction_loss	10.150286	10.221409

```
In [19]: fig, ax = plt.subplots(1,1,figsize=(14,8))
        diagnostic = pd.concat([pd.DataFrame(d.stack(), columns=[i]).T for i,d in enumerate(d
        cols = list('rgbcmyk')
        colors = defaultdict(lambda: cols.pop())
        for c in diagnostic:
            component, dataset = c
            kw = {}
            if dataset=='valid':
                kw['label'] = component
            else:
                kw['ls'] = '--'
            ax.plot(diagnostic[c].values, c=colors[component], **kw)
        ax.legend(bbox_to_anchor=(1, 0.7))
        ax.grid(True)
```



```
In [20]: model.eval()
pass
```

```
In [21]: # samples
z, x = model(128, mode='sample')
fig, ax = plt.subplots(1,1,figsize=(16,12))
ax.imshow(make_grid(
    x.data, nrow=16, range=(-1,1), normalize=True
).cpu().numpy().transpose(1,2,0), interpolation='nearest')
pass
```




```
In [22]: fig, ax = plt.subplots(1,1,figsize=(16,4))
         # training reconstructions
         x = cifar[:16]
         z, x_rec = model(x)
         ax.imshow(make_grid(
             torch.cat((x, x_rec.cpu().data)), nrow=16, range=(-1,1), normalize=True
         ).cpu().numpy().transpose(1,2,0), interpolation='nearest')
         pass
```



```
In [23]: fig, ax = plt.subplots(1,1,figsize=(16,4))
         # test reconstructions
         x = cifar_test[:16]
         z, x_rec = model(x)
         ax.imshow(make_grid(
             torch.cat((x, x_rec.cpu().data)), nrow=16, range=(-1,1), normalize=True
         ).cpu().numpy().transpose(1,2,0), interpolation='nearest')
         pass
```



pairwise distances give a sense of how encoded z are distributed compared to samples from the prior:

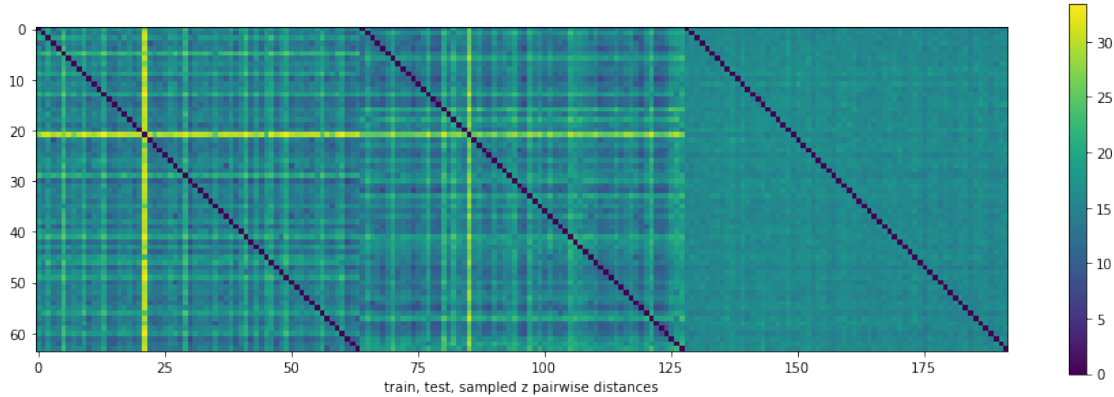
```
In [24]: fig, ax = plt.subplots(1,1,figsize=(16,5))
         n = 64

         import scipy.spatial
         train_z = model(cifar[:n], mode='encode').data.cpu().numpy()
         test_z = model(cifar_test[:n], mode='encode').data.cpu().numpy()
         sampled_z = model.sample_prior(n).data.cpu().numpy()
         train_cdist = scipy.spatial.distance.cdist(train_z,train_z)
         test_cdist = scipy.spatial.distance.cdist(test_z,test_z)
         sampled_cdist = scipy.spatial.distance.cdist(sampled_z,sampled_z)
```

```

cax = ax.imshow(np.concatenate((train_cdist, test_cdist, sampled_cdist),1))
fig.colorbar(cax)
ax.set_xlabel('train, test, sampled z pairwise distances')
pass

```



0.0.1 Interpolation

```

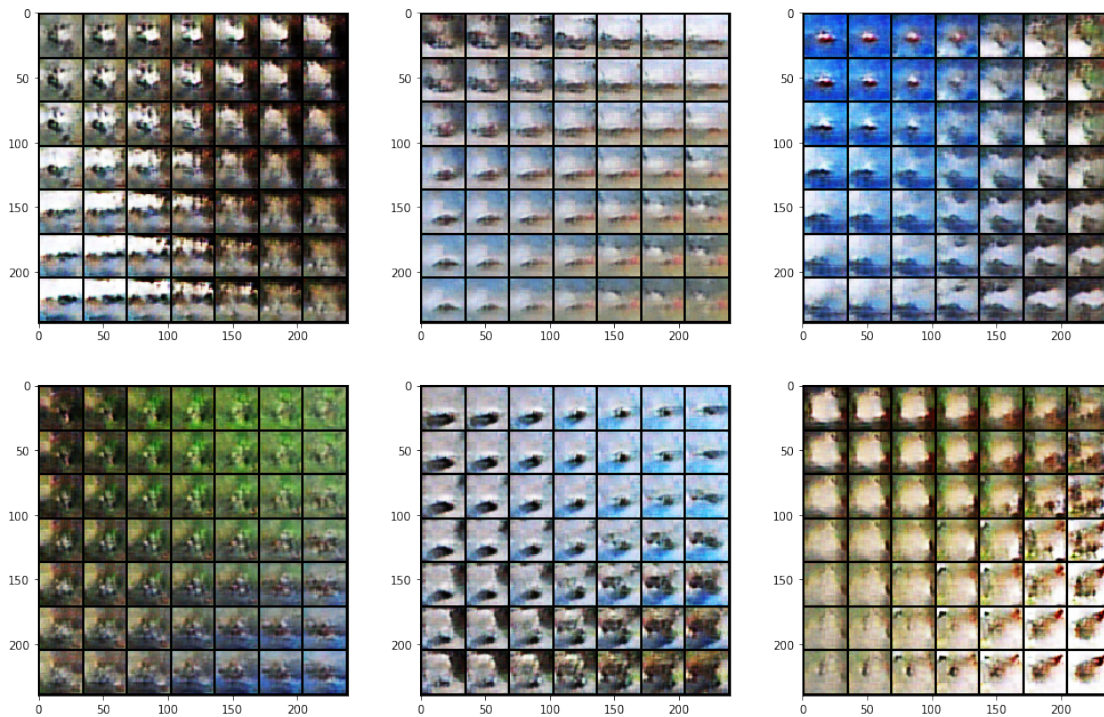
In [25]: def slerp(z0, z1, t):
    """Spherical linear interpolation over last dimension:
    z0.shape = z1.shape = (m,...,n, d) and t.shape = (q,) -> (q, m,...,n, d)
    Project to unit sphere and linearly interpolate magnitudes.
    """
    m0, m1 = (np.linalg.norm(z, 2, -1) for z in (z0, z1))
    p0, p1 = z0/np.expand_dims(m0,-1), z1/np.expand_dims(m1,-1)
    omega = np.arccos((p0*p1).sum(-1))
    while t.ndim<=omega.ndim:
        t = np.expand_dims(t,-1)
    sin_omega = np.sin(omega)
    t0 = np.sin((1-t)*omega)/sin_omega
    t1 = np.sin(t*omega)/sin_omega
    lim = abs(omega)<1e-15
    t1[lim] = t[lim]
    t0[lim] = (1-t)[lim]
    t0, t1 = np.expand_dims(t0, -1), np.expand_dims(t1, -1)
    m = np.expand_dims((1-t)*m0 + t*m1, -1)
    return m*(t0*p0 + t1*p1)
def slerp4(z, t):
    # z (4,n) and t (m,) -> z' (m, m, n)
    return slerp(
        slerp(z[0], z[1], t),
        slerp(z[2], z[3], t),
        t)

```

```

In [27]: n = 7
fig,ax = plt.subplots(2,3,figsize=(17,11))
for ax in ax.flatten():
    x = cifar_test[torch.LongTensor(np.random.randint(len(cifar_test),size=4))]
    z = model(x, mode='encode')
    z = slerp4(z.data.cpu().numpy(), np.linspace(0,1,n)).reshape(n*n,-1)
    x_rec = model(z, mode='generate')
    ax.imshow(make_grid(
        x_rec.cpu().data, nrow=n, range=(-1,1), normalize=True
    ).numpy().transpose(1,2,0), interpolation='nearest')

```



```

In [ ]:

```