

CIFAR

December 4, 2018

```
In [1]: use_gpu = True
        data_dir = '/home/kartik/data'

In [2]: import sys
        sys.path.append('../')

        %load_ext autoreload
        %autoreload 1
        %import alphagan

In [3]: from collections import defaultdict
        from psutil import cpu_count

        import numpy as np
        import pandas as pd

        import torch
        from torch import nn
        from torch.nn import init, Parameter
        import torch.nn.functional as F
        from torch.autograd import Variable
        from torch.utils.data import DataLoader
        from torchvision import datasets, models, transforms
        from torchvision.utils import make_grid

        import matplotlib.pyplot as plt
        %matplotlib inline

In [4]: # get a image data set into a tensor, scale to (-1, 1), discarding labels
        def torchvision_dataset(dset_fn, train=True):
            dset = dset_fn(
                data_dir,
                train=train,
                transform=transforms.Compose([
                    transforms.ToTensor()
                ]),
                target_transform=None,
                download=True)
            return torch.stack(list(zip(*dset))[0])*2-1
```

```

In [5]: cifar = torchvision_dataset(datasets.CIFAR10, train=True)
        cifar_test = torchvision_dataset(datasets.CIFAR10, train=False)
        print(cifar.size())

Files already downloaded and verified
Files already downloaded and verified
torch.Size([50000, 3, 32, 32])

In [6]: batch_size = 256

In [7]: num_workers = cpu_count() if use_gpu else 0

X_train = DataLoader(cifar, batch_size=batch_size, shuffle=True,
                      num_workers=num_workers, pin_memory=use_gpu)
X_test = DataLoader(cifar_test, batch_size=batch_size, shuffle=False,
                    num_workers=num_workers, pin_memory=use_gpu)

In [8]: class ChannelsToLinear(nn.Linear):
        """Flatten a Variable to 2d and apply Linear layer"""
        def forward(self, x):
            b = x.size(0)
            return super().forward(x.view(b,-1))

class LinearToChannels2d(nn.Linear):
        """Reshape 2d Variable to 4d after Linear layer"""
        def __init__(self, m, n, w=1, h=None, **kw):
            h = h or w
            super().__init__(m, n*w*h, **kw)
            self.w = w
            self.h = h
        def forward(self, x):
            b = x.size(0)
            return super().forward(x).view(b, -1, self.w, self.h)

class ResBlock(nn.Module):
        """Simple ResNet block"""
        def __init__(self, c,
                      activation=nn.LeakyReLU, norm=nn.BatchNorm2d,
                      init_gain=1, groups=1):
            super().__init__()
            self.a1 = activation()
            self.a2 = activation()
            self.norm1 = norm and norm(c)
            self.norm2 = norm and norm(c)

            to_init = []
            self.conv1 = nn.Conv2d(
                c, c, 3, 1, 1, bias=bool(norm), groups=groups)

```

```

to_init.append(self.conv1.weight)
self.conv2 = nn.Conv2d(
    c, c, 3, 1, 1, bias=bool(norm), groups=groups)
to_init.append(self.conv2.weight)

# if using grouping, add a 1x1 convolution to each conv layer
if groups!=1:
    self.conv1 = nn.Sequential(
        self.conv1, nn.Conv2d(c,c,1,bias=bool(norm)))
    self.conv2 = nn.Sequential(
        self.conv2, nn.Conv2d(c,c,1,bias=bool(norm)))
    to_init.extend([self.conv1[1].weight, self.conv2[1].weight])

# init
for w in to_init:
    init.xavier_normal(w, init_gain)

def forward(self, x):
    y = self.conv1(x)
    if self.norm1:
        y = self.norm1(y)
    y = self.a1(y)

    y = self.conv2(y)
    if self.norm2:
        y = self.norm2(y)

    return self.a2(x+y)

```

In [9]: latent_dim = 128

In [10]: *# encoder network*

```

h = 128
resample = nn.AvgPool2d
norm = nn.BatchNorm2d#None
a, g = nn.ReLU, init.calculate_gain('relu')
groups = 1#h//8
E = nn.Sequential(
    nn.Conv2d(3,h,5,1,2), resample(2), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    ChannelsToLinear(h*16, latent_dim)
)
for layer in (0,8):
    init.xavier_normal(E[layer].weight, g)

t = Variable(torch.randn(batch_size,3,32,32))
assert E(t).size() == (batch_size,latent_dim)

```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:15: UserWarning: nn.init.xavier_n
from ipykernel import kernelapp as app
```

```
In [11]: # generator network
```

```
h = 128
norm = nn.BatchNorm2d#None
a, g = nn.ReLU, init.calculate_gain('relu')
groups = 1#h//8
resample = lambda x: nn.Upsample(scale_factor=x)
G = nn.Sequential(
    LinearToChannels2d(latent_dim,h,4,4), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    nn.Conv2d(h, 3, 1), nn.Tanh()
)
for layer in (0,9):
    init.xavier_normal(G[layer].weight, g)

t = Variable(torch.randn(batch_size,latent_dim))
assert G(t).size() == (batch_size,3,32,32)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:16: UserWarning: nn.init.xavier_n
app.launch_new_instance()
```

```
In [12]: # discriminator network
```

```
h = 128
resample = nn.AvgPool2d
norm = nn.BatchNorm2d
a, g = lambda: nn.LeakyReLU(.2), init.calculate_gain('leaky_relu', .2)
groups = 1

D = nn.Sequential(
    nn.Conv2d(3,h,5,1,2), resample(2), a(),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups), resample(2),
    ResBlock(h, activation=a, norm=norm, init_gain=g, groups=groups),
    ChannelsToLinear(h*16, 1), nn.Sigmoid()
)
for layer in (0,8):
    init.xavier_normal(D[layer].weight, g)

t = Variable(torch.randn(batch_size,3,32,32))
assert D(t).size() == (batch_size,1)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:47: UserWarning: nn.init.xavier_n
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:16: UserWarning: nn.init.xavier_n
app.launch_new_instance()
```

```
In [13]: # code discriminator network
# no batch norm in the code discriminator, it causes trouble
h = 700
a, g = lambda: nn.LeakyReLU(.2), init.calculate_gain('leaky_relu', .2)
C = nn.Sequential(
    nn.Linear(latent_dim, h), a(),
    nn.Linear(h, h), a(),
    nn.Linear(h, 1), nn.Sigmoid(),
)

for i,layer in enumerate(C):
    if i%2==0:
        init.xavier_normal(layer.weight, g)

t = Variable(torch.randn(batch_size,latent_dim))
assert C(t).size() == (batch_size,1)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:13: UserWarning: nn.init.xavier_n
del sys.path[0]
```

```
In [14]: model = alphagan.AlphaGAN(E, G, D, C, latent_dim, lambd=40, z_lambd=0)
if use_gpu:
    model = model.cuda()
```

```
In [15]: # model.load_state_dict(torch.load('../models/cifar10_dim_{}_lambda_{}_zlambda_{}_epoch{}_
#         latent_dim, model.lambd, model.z_lambd, 8
#         )))
```

```
In [16]: diag = []
def log_fn(d):
    d = pd.DataFrame(d)
    diag.append(d)
    print(d)
def checkpoint_fn(model, epoch):
    path = '../models/cifar10_dim_{}_lambda_{}_zlambda_{}_epochs_{}.torch'.format(
        model.latent_dim, model.lambd, model.z_lambd, epoch
    )
    torch.save(model.state_dict(), path)
```

```
In [17]: model.fit(
    X_train, X_test,
    n_iter=(2,1,1), n_epochs=10,
    log_fn=log_fn, log_every=1,
```

```

        checkpoint_fn=checkpoint_fn, checkpoint_every=2
    )

HBox(children=(IntProgress(value=0, description='epoch', max=10, style=ProgressStyle(description='epoch', max=10, value=0)),

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle(description='training batch', max=196, value=0)),

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle(description='validating batch', max=40, value=0)),



|                         | train     | valid     |
|-------------------------|-----------|-----------|
| adversarial_loss        | 1.491249  | 1.208733  |
| code_adversarial_loss   | 1.107558  | 1.244499  |
| code_discriminator_loss | 2.846608  | 4.656633  |
| discriminator_loss      | 1.254535  | 1.358313  |
| reconstruction_loss     | 10.599710 | 12.200234 |



HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle(description='training batch', max=196, value=0)),

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle(description='validating batch', max=40, value=0)),



|                         | train    | valid    |
|-------------------------|----------|----------|
| adversarial_loss        | 1.158736 | 0.987269 |
| code_adversarial_loss   | 1.146322 | 1.761189 |
| code_discriminator_loss | 2.236347 | 1.193042 |
| discriminator_loss      | 1.314758 | 1.382846 |
| reconstruction_loss     | 8.604012 | 7.813487 |



HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle(description='training batch', max=196, value=0)),

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle(description='validating batch', max=40, value=0)),



|                         | train    | valid    |
|-------------------------|----------|----------|
| adversarial_loss        | 1.090680 | 0.892696 |
| code_adversarial_loss   | 1.134639 | 1.479062 |
| code_discriminator_loss | 2.156851 | 1.141951 |
| discriminator_loss      | 1.319052 | 1.321063 |
| reconstruction_loss     | 7.808319 | 7.550529 |



HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle(description='training batch', max=196, value=0)),

```

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

	train	valid
adversarial_loss	1.045439	0.839883
code_adversarial_loss	1.074357	1.322083
code_discriminator_loss	2.034314	1.191694
discriminator_loss	1.326508	1.404097
reconstruction_loss	7.283897	6.954408

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

	train	valid
adversarial_loss	1.020808	0.846527
code_adversarial_loss	1.013538	1.313122
code_discriminator_loss	1.769686	1.201549
discriminator_loss	1.330537	1.362522
reconstruction_loss	6.821696	6.706117

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

	train	valid
adversarial_loss	1.025753	0.811732
code_adversarial_loss	0.987070	1.252741
code_discriminator_loss	1.643611	1.132527
discriminator_loss	1.319158	1.366680
reconstruction_loss	6.568081	6.758773

HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle

HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle

	train	valid
adversarial_loss	1.043700	0.950884
code_adversarial_loss	0.934302	1.318607
code_discriminator_loss	1.592984	1.189724
discriminator_loss	1.305093	1.400025
reconstruction_loss	6.348385	6.494437

```
HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle
```

```
HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle
```

	train	valid
adversarial_loss	1.053522	0.962816
code_adversarial_loss	0.933363	1.177023
code_discriminator_loss	1.611025	1.205247
discriminator_loss	1.306511	1.387811
reconstruction_loss	6.211298	6.233727

```
HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle
```

```
HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle
```

	train	valid
adversarial_loss	1.066927	1.013610
code_adversarial_loss	0.916286	0.969998
code_discriminator_loss	1.553345	1.175605
discriminator_loss	1.307545	1.470947
reconstruction_loss	6.042090	6.057661

```
HBox(children=(IntProgress(value=0, description='training batch', max=196, style=ProgressStyle
```

```
HBox(children=(IntProgress(value=0, description='validating batch', max=40, style=ProgressStyle
```

	train	valid
adversarial_loss	1.060852	0.844432
code_adversarial_loss	0.902628	1.128680
code_discriminator_loss	1.530992	1.218899
discriminator_loss	1.311160	1.371434
reconstruction_loss	5.903250	6.049176

```
In [18]: fig, ax = plt.subplots(1,1,figsize=(14,8))
        diagnostic = pd.concat([pd.DataFrame(d.stack(), columns=[i]).T for i,d in enumerate(d
        cols = list('rgbcm')
        colors = defaultdict(lambda: cols.pop())
        for c in diagnostic:
            component, dataset = c
```




```

kw = {}
if dataset=='valid':
    kw['label'] = component
else:
    kw['ls'] = '--'
    ax.plot(diagnostic[c].values, c=colors[component], **kw)
ax.legend(bbox_to_anchor=(1, 0.7))

```

Out[18]: <matplotlib.legend.Legend at 0x7f33b345deb8>



```

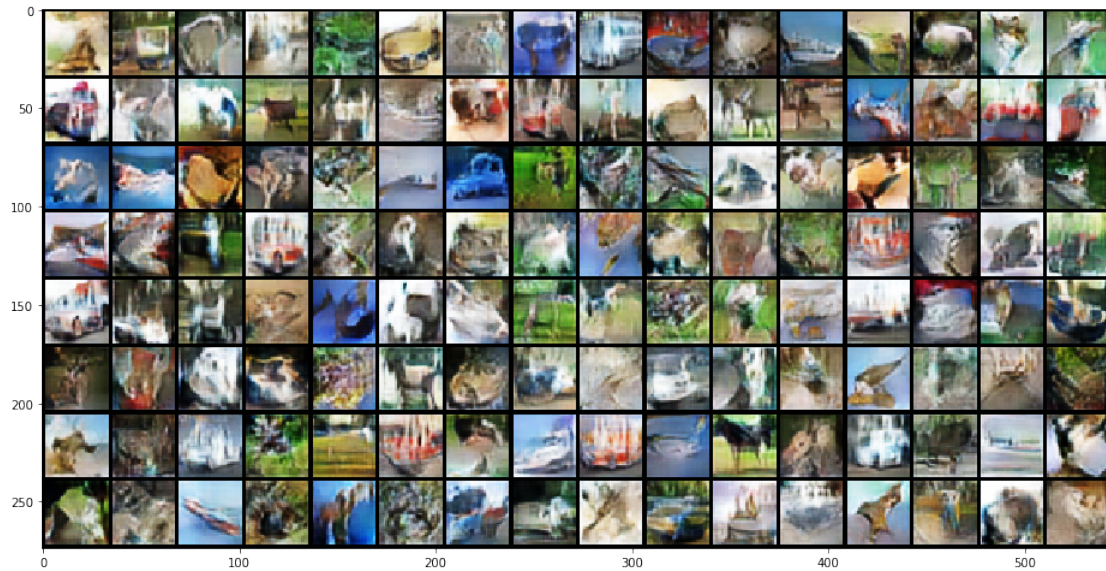
In [19]: model.eval()
pass

```


```

In [20]: # samples
z, x = model(128, mode='sample')
fig, ax = plt.subplots(1,1,figsize=(16,12))
ax.imshow(make_grid(
    x.data, nrow=16, range=(-1,1), normalize=True
).cpu().numpy().transpose(1,2,0), interpolation='nearest')
pass

```



```
In [21]: fig, ax = plt.subplots(1,1,figsize=(16,4))
         # training reconstructions
         x = cifar[:16]
         z, x_rec = model(x)
         ax.imshow(make_grid(
             torch.cat((x, x_rec.cpu().data)), nrow=16, range=(-1,1), normalize=True
         ).cpu().numpy().transpose(1,2,0), interpolation='nearest')
         pass
```



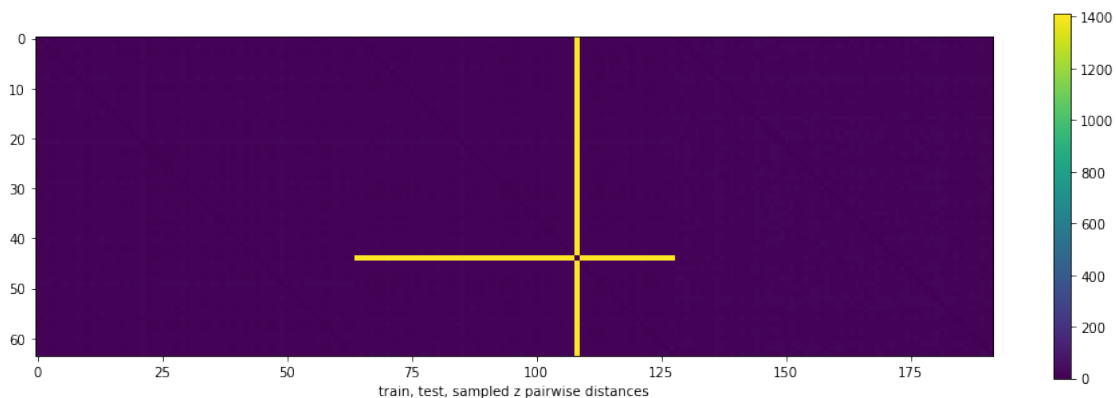
```
In [22]: fig, ax = plt.subplots(1,1,figsize=(16,4))
         # test reconstructions
         x = cifar_test[:16]
         z, x_rec = model(x)
         ax.imshow(make_grid(
             torch.cat((x, x_rec.cpu().data)), nrow=16, range=(-1,1), normalize=True
         ).cpu().numpy().transpose(1,2,0), interpolation='nearest')
         pass
```



0.0.1 pairwise distances give a sense of how encoded z are distributed compared the prior:

```
In [23]: fig, ax = plt.subplots(1,1,figsize=(16,5))
         n = 64

import scipy.spatial
train_z = model(cifar[:n], mode='encode').data.cpu().numpy()
test_z = model(cifar_test[:n], mode='encode').data.cpu().numpy()
sampled_z = model.sample_prior(n).data.cpu().numpy()
train_cdist = scipy.spatial.distance.cdist(train_z,train_z)
test_cdist = scipy.spatial.distance.cdist(test_z,test_z)
sampled_cdist = scipy.spatial.distance.cdist(sampled_z,sampled_z)
cax = ax.imshow(np.concatenate((train_cdist, test_cdist, sampled_cdist),1))
fig.colorbar(cax)
ax.set_xlabel('train, test, sampled z pairwise distances')
pass
```



0.0.2 interpolation

```
In [24]: def slerp(z0, z1, t):
         """Spherical linear interpolation over last dimension:
         z0.shape = z1.shape = (m,...n, d) and t.shape = (q,) -> (q, m,...n, d)
         Project to unit sphere and linearly interpolate magnitudes.
         """
```

```

m0, m1 = (np.linalg.norm(z, 2, -1) for z in (z0, z1))
p0, p1 = z0/np.expand_dims(m0,-1), z1/np.expand_dims(m1,-1)
omega = np.arccos((p0*p1).sum(-1))
while t.ndim<=omega.ndim:
    t = np.expand_dims(t,-1)
sin_omega = np.sin(omega)
t0 = np.sin((1-t)*omega)/sin_omega
t1 = np.sin(t*omega)/sin_omega
lim = abs(omega)<1e-15
t1[lim] = t[lim]
t0[lim] = (1-t)[lim]
t0, t1 = np.expand_dims(t0, -1), np.expand_dims(t1, -1)
m = np.expand_dims((1-t)*m0 + t*m1, -1)
return m*(t0*p0 + t1*p1)
def slerp4(z, t):
    # z (4,n) and t (m,) -> z' (m, m, n)
    return slerp(
        slerp(z[0], z[1], t),
        slerp(z[2], z[3], t),
        t)

```

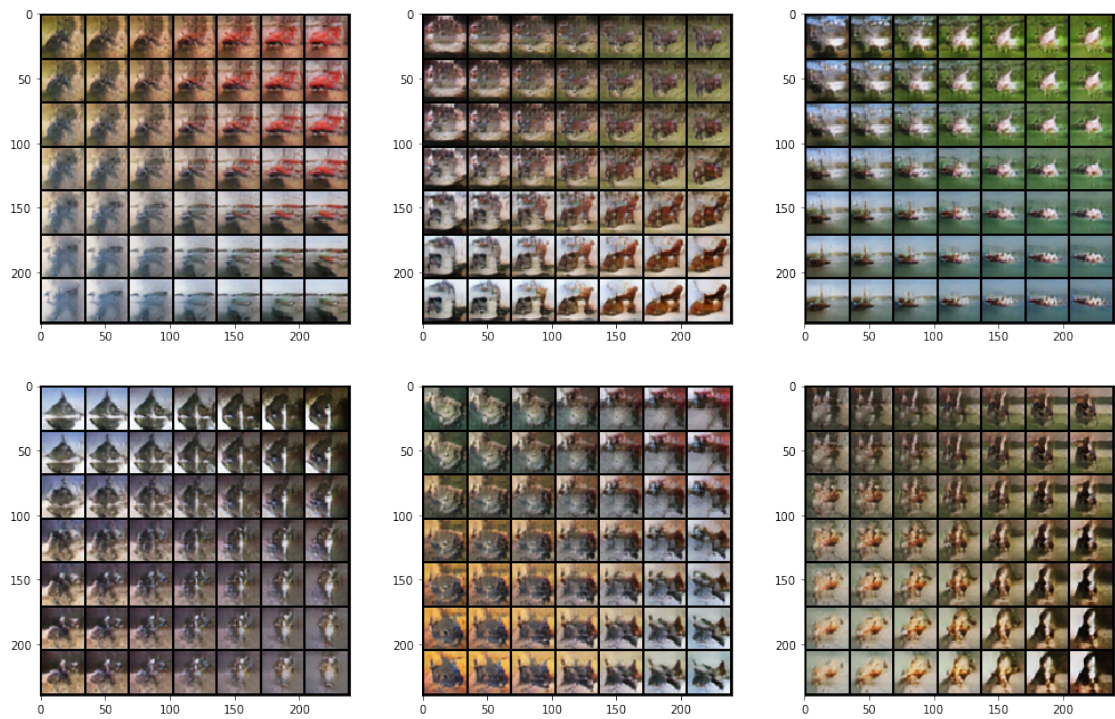
In [25]: # torch.manual_seed(2)

In [26]: n = 7

```

fig,ax = plt.subplots(2,3,figsize=(17,11))
for ax in ax.flatten():
    x = cifar_test[torch.LongTensor(np.random.randint(len(cifar_test),size=4))]
    z = model(x, mode='encode')
    z = slerp4(z.data.cpu().numpy(), np.linspace(0,1,n)).reshape(n*n,-1)
    x_rec = model(z, mode='generate')
    ax.imshow(make_grid(
        x_rec.cpu().data, nrow=n, range=(-1,1), normalize=True
    ).numpy().transpose(1,2,0), interpolation='nearest')

```



In []: